

Hall of Fame Expectancy

Daniel Antantis

Read in raw player data

```
hit_data = read.csv("~/Documents/MLB Project/allhitters.csv")
head(hit_data)
```

```
##      X      name      url Year Age      Tm Lg   G  PA  AB   R
## 1 0 Henry Aaron /players/a/aaronha01.shtml 1952  18 BSN-min C   87 345 345  NA
## 2 1 Henry Aaron /players/a/aaronha01.shtml 1953  19 MLN-min A  137 574 574  NA
## 3 2 Henry Aaron /players/a/aaronha01.shtml 1954  20      MLN NL  122 509 468  58
## 4 3 Henry Aaron /players/a/aaronha01.shtml 1955  21      MLN NL  153 665 602 105
## 5 4 Henry Aaron /players/a/aaronha01.shtml 1956  22      MLN NL  153 660 609 106
## 6 5 Henry Aaron /players/a/aaronha01.shtml 1957  23      MLN NL  151 675 615 118
##      H X2B X3B HR  RBI SB  CS  BB  SO    BA   OBP   SLG   OPS OPS.  TB  GDP HBP  SH  SF
## 1 116  19   4   9   NA  NA  NA  NA  NA 0.336 0.336 0.493 0.829   NA 170   NA   NA  NA  NA
## 2 208  36  14  22   NA  NA  NA  NA  NA 0.362 0.362 0.589 0.951   NA 338   NA   NA  NA  NA
## 3 131  27   6  13   69   2   2  28  39 0.280 0.322 0.447 0.769  104 209  13   3   6   4
## 4 189  37   9  27  106   3   1  49  61 0.314 0.366 0.540 0.906  141 325  20   3   7   4
## 5 200  34  14  26   92   2   4  37  54 0.328 0.365 0.558 0.923  151 340  21   2   5   7
## 6 198  27   6  44  132   1   1  57  58 0.322 0.378 0.600 0.978  166 369  13   0   0   3
##      IBB      Pos      Awards
## 1   NA      EAU  · NORL
## 2   NA      JCK  · SALL
## 3    0  *79/H      RoY-4
## 4    5  *974/H    AS,MVP-9
## 5    6   *9/H    AS,MVP-3
## 6   15  *98/H    AS,MVP-1
```

Read in Hall of Fame data and filter for only hitters by at-bats (AB)

```
url = "https://www.baseball-reference.com/awards/hof_batting.shtml"
tbl = url %>%
  read_html() %>%
  html_nodes('table') %>%
  html_table()
hof = data.frame(tbl)
hof = hof %>% filter(AB > 3000)
```

Filter for players eligible for HOF

```
hof_eligible = hit_data %>%
  na.omit() %>%
  group_by(name) %>%
  filter(AB > 200) %>%
```

```

summarise(
  Retirement_Year = max(Year),
  AB = sum(AB),
  HR = sum(HR),
  RBI = sum(RBI),
  H = sum(H),
  BA = mean(BA),
  OBP = mean(OBP),
  SB = sum(SB),
  BB = sum(BB),
  IBB = sum(IBB),
  SLG = mean(SLG),
  OPS = mean(OPS),
  TB = sum(TB),
  MVP = sum(ifelse(grepl("MVP-1", Awards, fixed = TRUE), 1, 0)),
  All_star = sum(ifelse(grepl("AS", Awards, fixed = TRUE), 1, 0)),
  Gold_glove = sum(ifelse(grepl("GG", Awards, fixed = TRUE), 1, 0)),
  Silver_slugger = sum(ifelse(grepl("SS", Awards, fixed = TRUE), 1, 0))
) %>%
filter(Retirement_Year < 2016)

```

Merge dataframes

```

hof_both = data.frame(hof_eligible$name[hof_eligible$name %in% hof$Name])
hof_both = hof_both %>%
  summarise(name = hof_eligible.name.hof_eligible.name..in..hof.Name.)
hof_both$hof = 1

```

Merge players with their stats (excluding players banned from HOF)

```

raw_data = hof_eligible %>%
  left_join(hof_both, by = "name")
raw_data[is.na(raw_data)] = 0
raw_data$ISO = raw_data$SLG - raw_data$BA
model_data = hof_eligible %>%
  left_join(hof_both, by = "name") %>%
  filter(!(name == "Barry Bonds" | name == "Pete Rose" | name == "Mark McGwire"))
model_data[is.na(model_data)] = 0
model_data$RC = ((model_data$H + model_data$BB) * model_data$TB) / (model_data$AB + model_data$BB)
model_data$ISO = model_data$SLG - model_data$BA

```

Split into training and testing sets

```
attach(model_data)
```

```

## The following object is masked _by_ .GlobalEnv:
##
##      hof

```

```

sample_size <- floor(0.8 * nrow(model_data))
set.seed(111)

```

```
train <- sample(seq_len(nrow(model_data)), size = sample_size)
model_train = model_data[train,]
model_test = model_data[-train,]
```

Linear Model

```
LDA.fit = lda(hof ~ HR + AB + BA + OPS + BB + MVP + All_star + Gold_glove + Silver_slugger, data = model_data)
LDA.pred = predict(LDA.fit, model_test)
mean(LDA.pred$class != model_test$hof)
```

```
## [1] 0.01976285
```

```
GLM.fit = glm(hof ~ HR + AB + BA + OPS + BB + ISO + MVP + All_star + Gold_glove + Silver_slugger, data = model_data)
GLM.probs = predict(GLM.fit, model_test, type = "response")
GLM.pred = rep(0, length(GLM.probs))
GLM.pred[GLM.probs > 0.5] = 1
mean(GLM.pred != model_test$hof)
```

```
## [1] 0.02371542
```

```
QDA.fit = qda(hof ~ HR + AB + BA + OPS + BB + MVP + All_star + Gold_glove + Silver_slugger, data = model_data)
QDA.pred = predict(QDA.fit, model_test)
mean(QDA.pred$class != model_test$hof)
```

```
## [1] 0.06126482
```

```
train.x = model.matrix(hof ~ HR + AB + BA + OPS + BB + MVP + All_star + Gold_glove + Silver_slugger, data = model_train)
train.hof = model_train$hof
test.x = model.matrix(hof ~ HR + AB + BA + OPS + BB + MVP + All_star + Gold_glove + Silver_slugger, data = model_test)
test.hof = model_test$hof

ridge.fit = cv.glmnet(train.x, train.hof, alpha = 0)
ridge.lambda = ridge.fit$lambda.min

ridge.probs = predict(ridge.fit, s = ridge.lambda, newx = test.x)
ridge.pred = rep(0, length(ridge.probs))
ridge.pred[ridge.probs > 0.4] = 1
mean(ridge.pred != model_test$hof)
```

```
## [1] 0.03162055
```

```
ridge.coef = predict(ridge.fit, type = "coefficients", s = ridge.lambda)
ridge.coef
```

```
## 11 x 1 sparse Matrix of class "dgCMatrix"
##              s1
## (Intercept) -1.481518e-02
## (Intercept) .
## HR          -5.883310e-05
```

```
## AB -9.889358e-06
## BA 2.791406e-02
## OPS 1.357075e-02
## BB 4.994942e-05
## MVP 1.127846e-01
## All_star 4.500345e-02
## Gold_glove 3.096033e-03
## Silver_slugger 1.795516e-02
```

Which players were classified incorrectly

```
incorrect_name = ifelse(GLM.pred != model_test$hof, model_test$name, NA)
incorrect_pred = ifelse(GLM.pred != model_test$hof, ridge.pred, NA)
incorrect = data.frame(incorrect_name, incorrect_pred)
incorrect = incorrect %>% na.omit()
```

Test model generalization

```
total.probs = predict(GLM.fit, model_data, type = "response")
total.pred = rep(0, length(total.probs))
total.pred[total.probs > 0.5] = 1
name = ifelse(total.pred != model_data$hof, model_data$name, NA)
pred = ifelse(total.pred != model_data$hof, total.pred, NA)
prob = ifelse(total.pred != model_data$hof, round(total.probs,3), NA)
end1 = data.frame(name, pred, prob)
end1 = end1 %>% na.omit()
end1_data = end1 %>% left_join(model_data, by = "name") %>% mutate_all(~replace(., is.na(.), 0))
head(end1_data)
```

```
##      name pred prob Retirement_Year  AB HR RBI  H  BA
## 1 Alan Trammell  0 0.107      1995  7950 183  976 2284 0.2848235
## 2 Bill Freehan  1 0.836      1976  6063 200  754 1587 0.2603571
## 3 Bill Mazeroski  0 0.134      1970  7498 137  834 1955 0.2578000
## 4 Billy Williams  0 0.238      1976  9270 424 1466 2693 0.2876875
## 5 Craig Biggio  0 0.275      2007 10753 288 1170 3034 0.2814211
## 6 Dale Murphy  1 0.627      1991  8278 418 1312 2192 0.2632500
##      OBP SB BB IBB SLG OPS TB MVP All_star Gold_glove
## 1 0.3508824 228 821 48 0.4141765 0.7651176 3344 0 6 4
## 2 0.3364286 24 625 67 0.4075714 0.7438571 2498 0 11 5
## 3 0.2992000 27 429 108 0.3638000 0.6630000 2775 0 7 8
## 4 0.3612500 90 1039 182 0.4870625 0.8481875 4569 0 6 0
## 5 0.3630000 408 1153 66 0.4322632 0.7950526 4668 0 7 4
## 6 0.3424375 170 1034 172 0.4625625 0.8051250 3881 2 7 5
##      Silver_slugger hof RC ISO
## 1 3 1 1183.8012 0.1293529
## 2 0 0 826.1926 0.1472143
## 3 0 1 834.5654 0.1060000
## 4 0 1 1654.0409 0.1993750
## 5 5 1 1641.6022 0.1508421
## 6 4 0 1344.5131 0.1993125
```

```
nrow(end1)
```

```
## [1] 44
```

```
name = ifelse(total.pred == 1 & model_data$hof == 1, model_data$name, NA)
pred = ifelse(total.pred == 1 & model_data$hof == 1, total.pred, NA)
prob = ifelse(total.pred == 1 & model_data$hof == 1, round(total.probs,3), NA)
end2 = data.frame(name, pred, prob)
end2 = end2 %>% na.omit()
end2_data = end2 %>% left_join(model_data, by = "name") %>% mutate_all(~replace(., is.na(.), 0))
head(end2_data)
```

```
##           name pred  prob Retirement_Year   AB  HR  RBI   H   BA
## 1      Al Kaline   1 0.989           1974 10088 398 1580 3000 0.2959048
## 2    Andre Dawson   1 0.665           1995  9784 436 1570 2738 0.2777895
## 3    Barry Larkin   1 0.922           2004  7622 193  924 2255 0.2969412
## 4 Brooks Robinson   1 0.993           1976 10424 264 1337 2801 0.2658421
## 5   Cal Ripken Jr.   1 1.000           2001 11512 431 1695 3179 0.2767500
## 6 Carl Yastrzemski   1 0.999           1983 11988 452 1844 3419 0.2838696
##           OBP  SB  BB  IBB      SLG      OPS   TB  MVP All_star Gold_glove
## 1 0.3747619 136 1276 133 0.4775714 0.8523333 4842   0    15         10
## 2 0.3218947 313  582 142 0.4815263 0.8033684 4737   1     8          8
## 3 0.3721176 368  903  63 0.4457059 0.8177647 3405   1    12          3
## 4 0.3196842  27  848 120 0.3954737 0.7152632 4197   1    15         16
## 5 0.3399000  36 1128 107 0.4501500 0.7900000 5163   2    19          2
## 6 0.3772609 168 1845 190 0.4580870 0.8355652 5539   1    18          7
## Silver_slugger hof      RC      ISO
## 1              0   1 1821.928 0.1816667
## 2              4   1 1517.156 0.2037368
## 3              9   1 1261.348 0.1487647
## 4              0   1 1358.663 0.1296316
## 5              8   1 1759.260 0.1734000
## 6              0   1 2107.807 0.1742174
```

```
nrow(end2)
```

```
## [1] 44
```

Lower Threshold

```
total.probs = predict(GLM.fit, model_data, type = "response")
total.pred = rep(0, length(total.probs))
total.pred[total.probs > 0.4] = 1
name = ifelse(total.pred != model_data$hof, model_data$name, NA)
pred = ifelse(total.pred != model_data$hof, total.pred, NA)
prob = ifelse(total.pred != model_data$hof, round(total.probs,3), NA)
end1 = data.frame(name, pred, prob)
end1 = end1 %>% na.omit()
end1_data = end1 %>% left_join(model_data, by = "name") %>% mutate_all(~replace(., is.na(.), 0))
head(end1_data)
```

```
##           name pred  prob Retirement_Year    AB  HR  RBI    H    BA
## 1  Alan Trammell    0 0.107             1995  7950 183  976 2284 0.2848235
## 2   Bill Freehan    1 0.836             1976  6063 200  754 1587 0.2603571
## 3 Bill Mazeroski    0 0.134             1970  7498 137  834 1955 0.2578000
## 4 Billy Williams    0 0.238             1976  9270 424 1466 2693 0.2876875
## 5   Craig Biggio    0 0.275             2007 10753 288 1170 3034 0.2814211
## 6   Dale Murphy    1 0.627             1991  8278 418 1312 2192 0.2632500
##           OBP  SB  BB  IBB      SLG      OPS  TB  MVP All_star Gold_glove
## 1 0.3508824 228  821  48 0.4141765 0.7651176 3344  0      6      4
## 2 0.3364286  24  625  67 0.4075714 0.7438571 2498  0     11      5
## 3 0.2992000  27  429 108 0.3638000 0.6630000 2775  0      7      8
## 4 0.3612500  90 1039 182 0.4870625 0.8481875 4569  0      6      0
## 5 0.3630000 408 1153  66 0.4322632 0.7950526 4668  0      7      4
## 6 0.3424375 170 1034 172 0.4625625 0.8051250 3881  2      7      5
## Silver_slugger hof      RC      ISO
## 1              3   1 1183.8012 0.1293529
## 2              0   0  826.1926 0.1472143
## 3              0   1  834.5654 0.1060000
## 4              0   1 1654.0409 0.1993750
## 5              5   1 1641.6022 0.1508421
## 6              4   0 1344.5131 0.1993125
```

```
nrow(end1)
```

```
## [1] 41
```

```
name = ifelse(total.pred == 1 & model_data$hof == 1, model_data$name, NA)
pred = ifelse(total.pred == 1 & model_data$hof == 1, total.pred, NA)
prob = ifelse(total.pred == 1 & model_data$hof == 1, round(total.probs,3), NA)
end2 = data.frame(name, pred, prob)
end2 = end2 %>% na.omit()
end2_data = end2 %>% left_join(model_data, by = "name") %>% mutate_all(~replace(., is.na(.), 0))
head(end2_data)
```

```
##           name pred  prob Retirement_Year    AB  HR  RBI    H    BA
## 1      Al Kaline    1 0.989             1974 10088 398 1580 3000 0.2959048
## 2   Andre Dawson    1 0.665             1995  9784 436 1570 2738 0.2777895
## 3   Barry Larkin    1 0.922             2004  7622 193  924 2255 0.2969412
## 4 Brooks Robinson    1 0.993             1976 10424 264 1337 2801 0.2658421
## 5   Cal Ripken Jr.    1 1.000             2001 11512 431 1695 3179 0.2767500
## 6 Carl Yastrzemski    1 0.999             1983 11988 452 1844 3419 0.2838696
##           OBP  SB  BB  IBB      SLG      OPS  TB  MVP All_star Gold_glove
## 1 0.3747619 136 1276 133 0.4775714 0.8523333 4842  0     15     10
## 2 0.3218947 313  582 142 0.4815263 0.8033684 4737  1      8      8
## 3 0.3721176 368  903  63 0.4457059 0.8177647 3405  1     12      3
## 4 0.3196842  27  848 120 0.3954737 0.7152632 4197  1     15     16
## 5 0.3399000  36 1128 107 0.4501500 0.7900000 5163  2     19      2
## 6 0.3772609 168 1845 190 0.4580870 0.8355652 5539  1     18      7
## Silver_slugger hof      RC      ISO
## 1              0   1 1821.928 0.1816667
## 2              4   1 1517.156 0.2037368
## 3              9   1 1261.348 0.1487647
## 4              0   1 1358.663 0.1296316
```

```
## 5          8    1 1759.260 0.1734000
## 6          0    1 2107.807 0.1742174
```

```
nrow(end2)
```

```
## [1] 49
```

GLM Deep Dive

```
GLM.fit$coefficients
```

```
##      (Intercept)          HR          AB          BA          OPS
## -1.614310e+01 -3.296638e-04  4.107136e-04 -9.011997e+01  5.560707e+01
##           BB          ISO          MVP        All_star        Gold_glove
## -3.677484e-03 -5.026194e+01  1.106218e+00  6.525447e-01 -7.060989e-03
## Silver_slugger
## -1.430774e-01
```

```
summary(GLM.fit)
```

```
##
## Call:
## glm(formula = hof ~ HR + AB + BA + OPS + BB + ISO + MVP + All_star +
##      Gold_glove + Silver_slugger, family = binomial, data = model_train,
##      subset = train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.1390  -0.1122  -0.0652  -0.0405   3.5985
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -1.614e+01  4.128e+00  -3.911 9.20e-05 ***
## HR           -3.297e-04  4.182e-03  -0.079  0.93717
## AB            4.107e-04  2.513e-04   1.635  0.10213
## BA           -9.012e+01  3.799e+01  -2.372  0.01767 *
## OPS           5.561e+01  1.973e+01   2.819  0.00482 **
## BB           -3.678e-03  1.638e-03  -2.244  0.02480 *
## ISO          -5.026e+01  2.382e+01  -2.110  0.03486 *
## MVP           1.106e+00  5.849e-01   1.891  0.05858 .
## All_star       6.525e-01  1.009e-01   6.465 1.01e-10 ***
## Gold_glove    -7.061e-03  9.888e-02  -0.071  0.94307
## Silver_slugger -1.431e-01  1.365e-01  -1.048  0.29452
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 466.22  on 1611  degrees of freedom
## Residual deviance: 181.48  on 1601  degrees of freedom
## (408 observations deleted due to missingness)
## AIC: 203.48
##
## Number of Fisher Scoring iterations: 9
```

Would the model incorrectly predict players who would be in HOF if not banned

```
banned = raw_data %>%
  filter(name == "Barry Bonds" | name == "Mark McGwire" | name == "Pete Rose")
head(banned)

## # A tibble: 3 x 20
##   name      Retir~1    AB    HR   RBI     H    BA   OBP    SB    BB   IBB   SLG
##   <chr>      <int> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 Barry Bon~    2007  9805   757  1986  2923  0.299  0.443   514  2549   685  0.613
## 2 Mark McGw~    2001  6281   596  1437  1658  0.264  0.394   13  1313   150  0.600
## 3 Pete Rose    1986 14331   160  1337  4328  0.297  0.371  199  1597   170  0.397
## # ... with 8 more variables: OPS <dbl>, TB <dbl>, MVP <dbl>, All_star <dbl>,
## #   Gold_glove <dbl>, Silver_slugger <dbl>, hof <dbl>, ISO <dbl>, and
## #   abbreviated variable name 1: Retirement_Year

GLM.probs = predict(GLM.fit, banned, type = "response")
GLM.pred = rep(0, length(GLM.probs))
GLM.pred[GLM.probs > 0.5] = 1
incorrect_ban_name = ifelse(GLM.pred != banned$hof, banned$name, NA)
incorrect_ban_pred = ifelse(GLM.pred != banned$hof, GLM.pred, NA)
incorrect_ban_prob = ifelse(GLM.pred != banned$hof, GLM.probs, NA)
incorrect_ban = data.frame(incorrect_ban_name, incorrect_ban_pred, incorrect_ban_prob)
incorrect_ban = incorrect_ban %>% na.omit()
head(incorrect_ban)

##   incorrect_ban_name incorrect_ban_pred incorrect_ban_prob
## 1      Barry Bonds                1      0.9999156
## 2      Mark McGwire                1      0.9684155
## 3      Pete Rose                  1      0.9964391
```

Would the model be better if only more modern players were considered?

```
modern = model_data %>%
  filter(Retirement_Year > 1960)

sample_size <- floor(0.8 * nrow(modern))
set.seed(111)
mod_train <- sample(seq_len(nrow(modern)), size = sample_size)
modern_train = modern[mod_train,]
modern_test = modern[-mod_train,]

GLM.modern.fit = glm(hof ~ HR + AB + BA + OPS + BB + ISO + MVP + All_star + Gold_glove + Silver_slugger
GLM.modern.probs = predict(GLM.modern.fit, model_test, type = "response")
GLM.modern.pred = rep(0, length(GLM.modern.probs))
GLM.modern.pred[GLM.modern.probs > 0.5] = 1
mean(GLM.modern.pred != modern_test$hof)

## Warning in GLM.modern.pred != modern_test$hof: longer object length is not a
## multiple of shorter object length

## [1] 0.0513834
```



```

modern.probs = predict(GLM.modern.fit, modern, type = "response")
modern.pred = rep(0, length(modern.probs))
modern.pred[modern.probs > 0.4] = 1
name = ifelse(modern.pred != modern$hof, modern$name, NA)
pred = ifelse(modern.pred != modern$hof, modern.pred, NA)
prob = ifelse(modern.pred != modern$hof, round(modern.probs,3), NA)
modern.end1 = data.frame(name, pred, prob)
modern.end1 = modern.end1 %>% na.omit()
modern.end1_data = modern.end1 %>% left_join(modern, by = "name") %>% mutate_all(~replace(., is.na(.), 0))
head(modern.end1_data)

```

```

##           name pred  prob Retirement_Year  AB  HR  RBI   H    BA
## 1      Al Oliver   1 0.491          1985 9380 219 1360 2843 0.3016667
## 2    Alan Trammell   0 0.096          1995 7950 183  976 2284 0.2848235
## 3    Bill Freehan   1 0.598          1976 6063 200  754 1587 0.2603571
## 4    Bill Mazeroski  0 0.045          1970 7498 137  834 1955 0.2578000
## 5    Billy Williams  0 0.321          1976 9270 424 1466 2693 0.2876875
## 6 Darryl Strawberry  1 0.573          1998 4905 315  925 1285 0.2611818
##           OBP  SB  BB  IBB      SLG      OPS  TB  MVP All_star Gold_glove
## 1 0.3424444  86  555 123 0.4421667 0.7846111 4206   0      7      0
## 2 0.3508824 228  821  48 0.4141765 0.7651176 3344   0      6      4
## 3 0.3364286  24  625  67 0.4075714 0.7438571 2498   0     11      5
## 4 0.2992000  27  429 108 0.3638000 0.6630000 2775   0      7      8
## 5 0.3612500  90 1039 182 0.4870625 0.8481875 4569   0      6      0
## 6 0.3579091 215  732 121 0.5161818 0.8742727 2535   0      8      0
## Silver_slugger hof      RC      ISO
## 1              3   0 1438.5494 0.1405000
## 2              3   1 1183.8012 0.1293529
## 3              0   0  826.1926 0.1472143
## 4              0   1  834.5654 0.1060000
## 5              0   1 1654.0409 0.1993750
## 6              2   0  907.0596 0.2550000

```

```
nrow(modern.end1)
```

```
## [1] 31
```

```

name = ifelse(modern.pred == 1 & modern$hof == 1, modern$name, NA)
pred = ifelse(modern.pred == 1 & modern$hof == 1, modern.pred, NA)
prob = ifelse(modern.pred == 1 & modern$hof == 1, round(modern.probs,3), NA)
modern.end2 = data.frame(name, pred, prob)
modern.end2 = modern.end2 %>% na.omit()
modern.end2_data = modern.end2 %>% left_join(modern, by = "name") %>% mutate_all(~replace(., is.na(.), 0))
head(modern.end2_data)

```

```

##           name pred  prob Retirement_Year  AB  HR  RBI   H    BA
## 1      Al Kaline   1 0.998          1974 10088 398 1580 3000 0.2959048
## 2    Andre Dawson   1 0.676          1995  9784 436 1570 2738 0.2777895
## 3    Barry Larkin   1 0.896          2004  7622 193  924 2255 0.2969412
## 4 Brooks Robinson   1 0.997          1976 10424 264 1337 2801 0.2658421
## 5   Cal Ripken Jr.   1 1.000          2001 11512 431 1695 3179 0.2767500
## 6 Carl Yastrzemski  1 1.000          1983 11988 452 1844 3419 0.2838696

```

```
##          OBP  SB   BB IBB          SLG          OPS   TB MVP All_star Gold_glove
## 1 0.3747619 136 1276 133 0.4775714 0.8523333 4842 0      15      10
## 2 0.3218947 313  582 142 0.4815263 0.8033684 4737 1       8       8
## 3 0.3721176 368  903  63 0.4457059 0.8177647 3405 1      12       3
## 4 0.3196842  27  848 120 0.3954737 0.7152632 4197 1      15      16
## 5 0.3399000  36 1128 107 0.4501500 0.7900000 5163 2      19       2
## 6 0.3772609 168 1845 190 0.4580870 0.8355652 5539 1      18       7
## Silver_slugger hof      RC      ISO
## 1              0   1 1821.928 0.1816667
## 2              4   1 1517.156 0.2037368
## 3              9   1 1261.348 0.1487647
## 4              0   1 1358.663 0.1296316
## 5              8   1 1759.260 0.1734000
## 6              0   1 2107.807 0.1742174
```

```
nrow(modern.end2)
```

```
## [1] 51
```

```
GLM.modern.fit$coefficients
```

```
##      (Intercept)          HR          AB          BA          OPS
## -3.221320e+01 -1.404267e-02  1.463007e-03 -1.441075e+02  9.002187e+01
##              BB          ISO          MVP          All_star  Gold_glove
## -4.979748e-03 -4.939107e+01  5.814346e-01  6.636787e-01  3.743445e-02
## Silver_slugger
## -1.905133e-01
```

```
summary(GLM.modern.fit)
```

```
##
## Call:
## glm(formula = hof ~ HR + AB + BA + OPS + BB + ISO + MVP + All_star +
##      Gold_glove + Silver_slugger, family = binomial, data = modern_train,
##      subset = train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.33889  -0.03234  -0.00824  -0.00215   2.95131
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -3.221e+01  7.400e+00  -4.353 1.34e-05 ***
## HR           -1.404e-02  6.790e-03  -2.068 0.038614 *
## AB            1.463e-03  4.314e-04   3.391 0.000695 ***
## BA           -1.441e+02  7.512e+01  -1.918 0.055061 .
## OPS           9.002e+01  4.091e+01   2.200 0.027773 *
## BB           -4.980e-03  2.905e-03  -1.714 0.086519 .
## ISO          -4.939e+01  4.871e+01  -1.014 0.310582
## MVP           5.814e-01  5.837e-01   0.996 0.319188
## All_star      6.637e-01  1.213e-01   5.471 4.47e-08 ***
## Gold_glove    3.743e-02  1.063e-01   0.352 0.724658
```

```
## Silver_slugger -1.905e-01  1.426e-01  -1.336 0.181510
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 361.17  on 1490  degrees of freedom
## Residual deviance: 102.70  on 1480  degrees of freedom
## (529 observations deleted due to missingness)
## AIC: 124.7
##
## Number of Fisher Scoring iterations: 10
```

Filter out players who have retired so that all that remains are players currently playing or not eligible for the HOF

```
current = hit_data %>%
  na.omit() %>%
  group_by(name) %>%
  filter(AB > 200) %>%
  summarise(
    Retirement_Year = max(Year),
    AB = sum(AB),
    HR = sum(HR),
    RBI = sum(RBI),
    H = sum(H),
    BA = mean(BA),
    OBP = mean(OBP),
    SB = sum(SB),
    BB = sum(BB),
    IBB = sum(IBB),
    SLG = mean(SLG),
    OPS = mean(OPS),
    TB = sum(TB),
    MVP = sum(ifelse(grepl("MVP-1,", Awards, fixed = TRUE), 1, 0)),
    All_star = sum(ifelse(grepl("AS", Awards, fixed = TRUE), 1, 0)),
    Gold_glove = sum(ifelse(grepl("GG", Awards, fixed = TRUE), 1, 0)),
    Silver_slugger = sum(ifelse(grepl("SS", Awards, fixed = TRUE), 1, 0))
  ) %>%
  filter(Retirement_Year > 2016)
hof_both = data.frame(current$name[current$name %in% hof$Name])
hof_both = hof_both %>%
  summarise(name = current.name.current.name..in..hof.Name.)
hof_both$hof = 1
raw_data = current %>%
  left_join(hof_both, by = "name")
raw_data[is.na(raw_data)] = 0
raw_data$ISO = raw_data$SLG - raw_data$BA
curr_data = current %>%
  left_join(hof_both, by = "name") %>%
  filter(!(name == "Barry Bonds" | name == "Pete Rose" | name == "Mark McGwire"))
curr_data[is.na(curr_data)] = 0
curr_data$RC = ((curr_data$H + curr_data$BB) * curr_data$TB) / (curr_data$AB + curr_data$BB)
curr_data$ISO = curr_data$SLG - curr_data$BA
```

Test on current players and players not yet inducted into HOF

```
curr.probs = predict(GLM.fit, curr_data, type = "response")
curr.pred = rep(0, length(curr.probs))
curr.pred[curr.probs > 0.4] = 1
name = ifelse(curr.pred == 1, curr_data$name, NA)
pred = ifelse(curr.pred == 1, curr.pred, NA)
prob = ifelse(curr.pred == 1, round(curr.probs,3), NA)
curr1 = data.frame(name, pred, prob)
curr1 = curr1 %>% na.omit()
curr1_data = curr1 %>% left_join(curr_data, by = "name") %>% mutate_all(~replace(., is.na(.), 0))
head(curr1_data)
```

```
##           name pred  prob Retirement_Year  AB  HR  RBI   H   BA
## 1  Albert Pujols   1 0.949             2019 10687 656 2075 3202 0.2985789
## 2   Bryce Harper   1 0.653             2021  4127 234  669 1140 0.2773333
## 3   Buster Posey   1 0.532             2021  4597 148  680 1398 0.3041000
## 4  Carlos Beltran   1 0.576             2017 11021 510 1814 3080 0.2783913
## 5 Freddie Freeman   1 0.467             2021  5472 258  907 1610 0.2956364
## 6    Joey Votto     1 0.400             2021  6222 291  967 1895 0.3008462
##           OBP  SB  BB  IBB      SLG      OPS   TB  MVP All_star Gold_glove
## 1 0.3754211 114 1322 311 0.5460000 0.9213158 5863   2    10         2
## 2 0.3820000  97  719  85 0.5103333 0.8924444 2113   1     6         0
## 3 0.3741000  20  497  63 0.4657000 0.8397000 2133   1     7         1
## 4 0.3526522 354 1255 123 0.4911739 0.8438261 5429   0     9         3
## 5 0.3873636  49  746  96 0.5159091 0.9032727 2785   1     5         1
## 6 0.4145385  79 1203 140 0.5102308 0.9249231 3217   0     6         1
## Silver_slugger hof      RC      ISO
## 1              6   0 2208.6945 0.2474211
## 2              1   0  810.5792 0.2330000
## 3              4   0  793.4894 0.1616000
## 4              2   0 1917.1322 0.2127826
## 5              2   0 1055.2364 0.2202727
## 6              0   0 1342.2580 0.2093846
```

```
nrow(curr1)
```

```
## [1] 11
```

Players who are on the edge

```
curr2.pred = rep(0, length(curr.probs))
curr2.pred[curr.probs > 0.25] = 1
name = ifelse(curr2.pred == 1, curr_data$name, NA)
pred = ifelse(curr2.pred == 1, curr.pred, NA)
prob = ifelse(curr2.pred == 1, round(curr.probs,3), NA)
curr2 = data.frame(name, pred, prob)
curr2 = curr2 %>% na.omit()
curr2_data = curr2 %>% left_join(curr_data, by = "name") %>% mutate_all(~replace(., is.na(.), 0)) %>% f
head(curr2_data)
```

```
##           name pred  prob Retirement_Year  AB  HR  RBI   H   BA
```

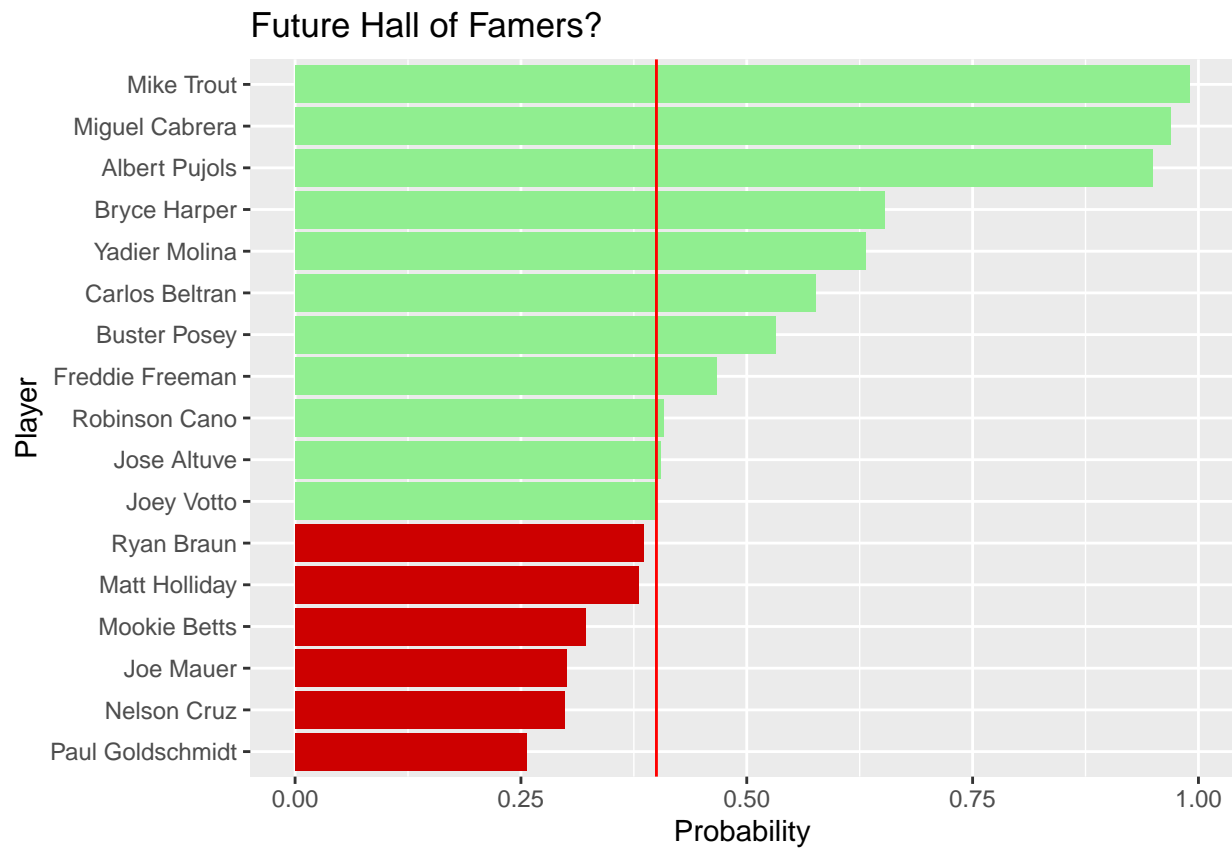
```
## 1      Joe Mauer      0 0.301      2018 6823 137 906 2090 0.3052857
## 2      Matt Holliday 0 0.380      2017 7537 338 1326 2263 0.2979375
## 3      Mookie Betts  0 0.322      2021 3538 163 529 1052 0.2945714
## 4      Nelson Cruz   0 0.298      2021 6235 406 1117 1736 0.2785385
## 5 Paul Goldschmidt  0 0.256      2021 4752 248 830 1387 0.2908889
## 6      Ryan Braun    0 0.386      2019 6493 344 1128 1933 0.2957692
##      OBP  SB  BB  IBB      SLG      OPS  TB MVP All_star Gold_glove
## 1 0.3870714 51 928 145 0.4340714 0.8210714 2979 1 6 3
## 2 0.3781250 122 862 65 0.5053125 0.8833125 3852 0 7 0
## 3 0.3727143 136 421 28 0.5242857 0.8970000 1846 1 5 5
## 4 0.3463846 74 614 65 0.5270000 0.8734615 3302 0 7 0
## 5 0.3862222 128 747 105 0.5162222 0.9022222 2469 0 6 3
## 6 0.3579231 215 579 56 0.5291538 0.8871538 3462 1 6 0
## Silver_slugger hof      RC      ISO
## 1      5 0 1159.9306 0.1287857
## 2      4 0 1433.2063 0.2073750
## 3      4 0 686.8295 0.2297143
## 4      3 0 1132.9683 0.2484615
## 5      4 0 958.1462 0.2253333
## 6      5 0 1229.7149 0.2333846
```

```
nrow(curr2_data)
```


```
## [1] 6
```

```
curr = curr1_data %>% rbind(curr2_data)
```

```
ggplot(data = curr, mapping = aes(x = reorder(name, prob), y = prob, fill = as.factor(pred))) + geom_bar
```



```
ggplot(data = modern.end1_data, mapping = aes(x = reorder(name, prob), y = prob, fill = as.factor(pred)))
```



```
ggplot(data = incorrect_ban, mapping = aes(x = reorder(incorrect_ban_name, incorrect_ban_prob), y = incorrect_ban_name))
```

